

## ▼ AlexNet Architecture

Following the tutorial

<https://www.kaggle.com/code/vortexkol/alexnet-cnn-architecture-on-tensorflow-beginner>

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import os
import time

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

+ 코드

+ 텍스트

## ▼ Dataset

```
(train_images, train_labels), (test_images, test_labels) = keras.datasets.cifar10.load_data()
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 4s 0us/step
```

In order to reference the class names of the images during the visualization stage, a python list containing the classes is initialized with the variable name CLASS\_NAMES.

```
CLASS_NAMES= ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

tf.data.Dataset.from\_tensor\_slices method takes the train, test, and validation dataset partitions and returns a corresponding TensorFlow Dataset representation.

```
train_ds=tf.data.Dataset.from_tensor_slices((train_images,train_labels))
test_ds=tf.data.Dataset.from_tensor_slices((test_images,test_labels))
```

## ▼ Preprocessing

Usually, preprocessing is conducted to ensure the data utilized is within an appropriate format.

First, let's visualize the images within the CIFAR-10 dataset.

The CIFAR-10 images have small dimensions, which makes visualization of the actual pictures a bit difficult.

```
plt.figure(figsize=(30,30))
for i,(image,label) in enumerate(train_ds.take(20)):
    #print(label)
    ax=plt.subplot(5,5,i+1)
    plt.imshow(image)
    plt.title(CLASS_NAMES[label.numpy()[0]])
    plt.axis('off')
```



## Preprocessing functions

```
import tensorflow as tf
import numpy as np
import pandas as pd

import random as python_random
from tensorflow.keras.utils import to_categorical
from sklearn.model_selection import train_test_split
from tensorflow.keras.datasets import cifar10

def zero_one_scaler(image):
    return image/255.0

def get_preprocessed_ohe(images, labels, pre_func=None):
    # preprocessing 함수가 입력되면 이를 이용하여 image array를 scaling 적용.
    if pre_func is not None:
        images = pre_func(images)
    # OHE 적용
    oh_labels = to_categorical(labels)
    return images, oh_labels

# 학습/검증/테스트 데이터 세트에 전처리 및 OHE 적용한 뒤 반환
def get_train_valid_test_set(train_images, train_labels, test_images, test_labels, valid_size=0.15, random_state=2021):
    # 학습 및 테스트 데이터 세트를 0 ~ 1사이값 float32로 변경 및 OHE 적용.
    train_images, train_oh_labels = get_preprocessed_ohe(train_images, train_labels)
    test_images, test_oh_labels = get_preprocessed_ohe(test_images, test_labels)

    # 학습 데이터를 검증 데이터 세트로 다시 분리
    tr_images, val_images, tr_oh_labels, val_oh_labels = train_test_split(train_images, train_oh_labels, test_size=valid_size, random_state=random_stat

    return (tr_images, tr_oh_labels), (val_images, val_oh_labels), (test_images, test_oh_labels)

# CIFAR10 데이터 재 로딩 및 Scaling/OHE 전처리 적용하여 학습/검증/데이터 세트 생성.
(train_images, train_labels), (test_images, test_labels) = cifar10.load_data()
print(train_images.shape, train_labels.shape, test_images.shape, test_labels.shape)
(tr_images, tr_oh_labels), (val_images, val_oh_labels), (test_images, test_oh_labels) = W
get_train_valid_test_set(train_images, train_labels, test_images, test_labels, valid_size=0.2, random_state=2021)

print(tr_images.shape, tr_oh_labels.shape, val_images.shape, val_oh_labels.shape, test_images.shape, test_oh_labels.shape)

(50000, 32, 32, 3) (50000, 1) (10000, 32, 32, 3) (10000, 1)
(40000, 32, 32, 3) (40000, 10) (10000, 32, 32, 3) (10000, 10) (10000, 32, 32, 3) (10000, 10)

IMAGE_SIZE = 227
BATCH_SIZE = 64

from tensorflow.keras.utils import Sequence
import cv2
import sklearn

# 입력 인자 images_array labels는 모두 numpy array로 들어옴.
# 인자로 입력되는 images_array는 전체 32x32 image array임.
class CIFAR_Dataset(Sequence):
    def __init__(self, images_array, labels, batch_size=BATCH_SIZE, augmentor=None, shuffle=False, pre_func=None):
        ...

    파라미터 설명
    images_array: 원본 32x32 만큼의 image 배열값.
    labels: 해당 image의 label들
    batch_size: __getitem__(self, index) 호출 시 마다 가져올 데이터 batch 건수
    augmentor: albumentations 객체
    shuffle: 학습 데이터의 경우 epoch 종료시마다 데이터를 섞을지 여부
```

```

'''
# 객체 생성 인자로 들어온 값을 객체 내부 변수로 할당.
# 인자로 입력되는 images_array는 전체 32x32 image array임.
self.images_array = images_array
self.labels = labels
self.batch_size = batch_size
self.augmentor = augmentor
self.pre_func = pre_func
# train data의 경우
self.shuffle = shuffle
if self.shuffle:
    # 객체 생성시에 한번 데이터를 섞음.
    #self.on_epoch_end()
    pass

# Sequence를 상속받은 Dataset은 batch_size 단위로 입력된 데이터를 처리함.
# __len__()은 전체 데이터 건수가 주어졌을 때 batch_size단위로 몇번 데이터를 반환하는지 나타냄
def __len__(self):
    # batch_size단위로 데이터를 몇번 가져와야하는지 계산하기 위해 전체 데이터 건수를 batch_size로 나누되, 정수로 정확히 나뉘지 않을 경우 1회를 더
    return int(np.ceil(len(self.labels) / self.batch_size))

# batch_size 단위로 image_array, label_array 데이터를 가져와서 변환한 뒤 다시 반환함
# 인자로 몇번째 batch 인지를 나타내는 index를 입력하면 해당 순서에 해당하는 batch_size 만큼의 데이터를 가공하여 반환
# batch_size 갯수만큼 변환된 image_array와 label_array 반환.
def __getitem__(self, index):
    # index는 몇번째 batch인지를 나타냄.
    # batch_size만큼 순차적으로 데이터를 가져오려면 array에서 index*self.batch_size:(index+1)*self.batch_size 만큼의 연속 데이터를 가져오면 됨
    # 32x32 image array를 self.batch_size만큼 가져옴.
    images_fetch = self.images_array[index*self.batch_size:(index+1)*self.batch_size]
    if self.labels is not None:
        label_batch = self.labels[index*self.batch_size:(index+1)*self.batch_size]

    # 만일 객체 생성 인자로 albumentation으로 만든 augmentor가 주어진다면 아래와 같이 augmentor를 이용하여 image 변환
    # augmentations은 개별 image만 변환할 수 있으므로 batch_size만큼 할당된 image_name_batch를 한 건씩 iteration하면서 변환 수행.
    # 변환된 image 배열값을 담은 image_batch 선언. image_batch 배열은 float32 로 설정.
    image_batch = np.zeros((images_fetch.shape[0], IMAGE_SIZE, IMAGE_SIZE, 3), dtype='float32')

    # batch_size에 담긴 건수만큼 iteration 하면서 opencv image load -> image augmentation 변환(augmentor가 not None일 경우)-> image_batch에 담음.
    for image_index in range(images_fetch.shape[0]):
        #image = cv2.cvtColor(cv2.imread(image_name_batch[image_index]), cv2.COLOR_BGR2RGB)
        # 원본 image를 IMAGE_SIZE x IMAGE_SIZE 크기로 변환
        image = cv2.resize(images_fetch[image_index], (IMAGE_SIZE, IMAGE_SIZE))
        # 만약 augmentor가 주어졌다면 이를 적용.
        if self.augmentor is not None:
            image = self.augmentor(image=image)['image']

        # 만약 scaling 함수가 입력되었다면 이를 적용하여 scaling 수행.
        if self.pre_func is not None:
            image = self.pre_func(image)

        # image_batch에 순차적으로 변환된 image를 담음.
        image_batch[image_index] = image

    return image_batch, label_batch

# epoch가 한번 수행이 완료 될 때마다 모델의 fit()에서 호출됨.
def on_epoch_end(self):
    if(self.shuffle):
        #print('epoch end')
        # 원본 image배열과 label을 쌍을 맞춰서 섞어준다. scikit learn의 utils.shuffle에서 해당 기능 제공
        self.images_array, self.labels = sklearn.utils.shuffle(self.images_array, self.labels)
    else:
        pass

def zero_one_scaler(image):
    return image/255.0

# CIFAR10 데이터 재 로딩 및 Scaling/OHE 전처리 적용하여 학습/검증/데이터 세트 생성.
(train_images, train_labels), (test_images, test_labels) = keras.datasets.cifar10.load_data()
print(train_images.shape, train_labels.shape, test_images.shape, test_labels.shape)

# Split data
(tr_images, tr_oh_labels), (val_images, val_oh_labels), (test_images, test_oh_labels) = W
get_train_valid_test_set(train_images, train_labels, test_images, test_labels, valid_size=0.2, random_state=2021)
print(tr_images.shape, tr_oh_labels.shape, val_images.shape, val_oh_labels.shape, test_images.shape, test_oh_labels.shape)

# Resize sequence dataset
tr_ds = CIFAR_Dataset(tr_images, tr_oh_labels, batch_size=BATCH_SIZE, augmentor=None, shuffle=True, pre_func=zero_one_scaler)

```

```
val_ds = CIFAR_Dataset(val_images, val_labels, batch_size=BATCH_SIZE, augmentor=None, shuffle=False, pre_func=zero_one_scaler)
```

```
print(next(iter(tr_ds))[0].shape, next(iter(val_ds))[0].shape)
print(next(iter(tr_ds))[1].shape, next(iter(val_ds))[1].shape)
print(next(iter(tr_ds))[0][0])
```

```
(50000, 32, 32, 3) (50000, 1) (10000, 32, 32, 3) (10000, 1)
(40000, 32, 32, 3) (40000, 10) (10000, 32, 32, 3) (10000, 10) (10000, 32, 32, 3) (10000, 10)
(64, 227, 227, 3) (64, 227, 227, 3)
(64, 10) (64, 10)
[[[0.6431373  0.68235296 0.69411767]
  [0.6431373  0.68235296 0.69411767]
  [0.6431373  0.68235296 0.69411767]
  ...
  [0.27450982 0.28235295 0.26666668]
  [0.27450982 0.28235295 0.26666668]
  [0.27450982 0.28235295 0.26666668]]

[[[0.6431373  0.68235296 0.69411767]
  [0.6431373  0.68235296 0.69411767]
  [0.6431373  0.68235296 0.69411767]
  ...
  [0.27450982 0.28235295 0.26666668]
  [0.27450982 0.28235295 0.26666668]
  [0.27450982 0.28235295 0.26666668]]

[[[0.6431373  0.68235296 0.69411767]
  [0.6431373  0.68235296 0.69411767]
  [0.6431373  0.68235296 0.69411767]
  ...
  [0.27450982 0.28235295 0.26666668]
  [0.27450982 0.28235295 0.26666668]
  [0.27450982 0.28235295 0.26666668]]

...

[[[0.91764706 0.8627451  0.88235295]
  [0.91764706 0.8627451  0.88235295]
  [0.91764706 0.8627451  0.88235295]
  ...
  [0.27058825 0.28235295 0.26666668]
  [0.27058825 0.28235295 0.26666668]
  [0.27058825 0.28235295 0.26666668]]

[[[0.91764706 0.8627451  0.88235295]
  [0.91764706 0.8627451  0.88235295]
  [0.91764706 0.8627451  0.88235295]
  ...
  [0.27058825 0.28235295 0.26666668]
  [0.27058825 0.28235295 0.26666668]
  [0.27058825 0.28235295 0.26666668]]

[[[0.91764706 0.8627451  0.88235295]
  [0.91764706 0.8627451  0.88235295]
  [0.91764706 0.8627451  0.88235295]
  ...
  [0.27058825 0.28235295 0.26666668]
  [0.27058825 0.28235295 0.26666668]
  [0.27058825 0.28235295 0.26666668]]]
```

## ▼ Data Pipeline

An input/data pipeline is described as a series of functions or methods that are called consecutively one after another.

Input pipelines are a chain of functions that either act upon the data or enforces an operation on the data flowing through the pipeline.

## ▼ AlexNet Model Implementation

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Dense, Conv2D, Dropout, Flatten, Activation, MaxPooling2D
from tensorflow.keras.optimizers import Adam, RMSprop
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint, LearningRateScheduler
from tensorflow.keras import regularizers

def create_alexnet(in_shape=(227, 227, 3), n_classes=10, kernel_regularizer=None):
    # 첫번째 CNN->ReLU->MaxPool, kernel_size를 매우 크게 가져감(11, 11). 지금은 (11, 11) 사용하지 않음.
    input_tensor = Input(shape=in_shape)

    x = Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), padding='valid')(input_tensor)
    x = Activation('relu')(x)
```

```

x = tf.nn.local_response_normalization(x)
x = MaxPooling2D(pool_size=(3,3), strides=(2,2))(x)

# 두번째 CNN->ReLU->MaxPool. kernel_size=(5, 5)
x = Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), padding='same', kernel_regularizer=kernel_regularizer)(x)
x = Activation('relu')(x)
x = tf.nn.local_response_normalization(x)
x = MaxPooling2D(pool_size=(3,3), strides=(2,2))(x)

# 3x3 Conv 2번 연속 적용. filters는 384개
x = Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same', kernel_regularizer=kernel_regularizer)(x)
x = Activation('relu')(x)
x = tf.nn.local_response_normalization(x)

x = Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), padding='same', kernel_regularizer=kernel_regularizer)(x)
x = Activation('relu')(x)
x = tf.nn.local_response_normalization(x)

# 3x3 Conv를 적용하되 filters 수를 줄이고 maxpooling을 적용
x = Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), padding='same', kernel_regularizer=kernel_regularizer)(x)
x = Activation('relu')(x)
x = BatchNormalization()(x)
x = MaxPooling2D(pool_size=(3,3), strides=(2,2))(x)

# Dense 연결을 위한 Flatten
x = Flatten()(x)

# Dense + Dropout을 연속 적용.
x = Dense(units = 4096, activation = 'relu')(x)
x = Dropout(0.5)(x)

x = Dense(units = 4096, activation = 'relu')(x)
x = Dropout(0.5)(x)

# 마지막 softmax 층 적용.
output = Dense(units = n_classes, activation = 'softmax')(x)

model = Model(inputs=input_tensor, outputs=output)
model.summary()

return model

```

```

# Create model with input 128x128x3
model = create_alexnet(in_shape=(227, 227, 3), n_classes=10, kernel_regularizer=regularizers.l2(l2=1e-4))

```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 227, 227, 3)]	0
conv2d (Conv2D)	(None, 55, 55, 96)	34944
activation (Activation)	(None, 55, 55, 96)	0
tf.nn.local_response_normalization (TFOpLambda)	(None, 55, 55, 96)	0
max_pooling2d (MaxPooling2D)	(None, 27, 27, 96)	0
conv2d_1 (Conv2D)	(None, 27, 27, 256)	614656
activation_1 (Activation)	(None, 27, 27, 256)	0
tf.nn.local_response_normalization_1 (TFOpLambda)	(None, 27, 27, 256)	0
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 256)	0
conv2d_2 (Conv2D)	(None, 13, 13, 384)	885120
activation_2 (Activation)	(None, 13, 13, 384)	0
tf.nn.local_response_normalization_2 (TFOpLambda)	(None, 13, 13, 384)	0
conv2d_3 (Conv2D)	(None, 13, 13, 384)	1327488
activation_3 (Activation)	(None, 13, 13, 384)	0
tf.nn.local_response_normalization_3 (TFOpLambda)	(None, 13, 13, 384)	0

conv2d_4 (Conv2D)	(None, 13, 13, 256)	884992
activation_4 (Activation)	(None, 13, 13, 256)	0
batch_normalization (Batch Normalization)	(None, 13, 13, 256)	1024
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 256)	0
flatten (Flatten)	(None, 9216)	0
dense (Dense)	(None, 4096)	37752832
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16781312


```
model.compile(optimizer=Adam(lr=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# 5번 iteration내에 validation loss가 향상되지 않으면 learning rate을 기존 learning rate * 0.2로 줄임.
r_lr_cb = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=5, mode='min', verbose=1)
e_ly_cb = EarlyStopping(monitor='val_loss', patience=10, mode='min', verbose=1)
```

```
history = model.fit(tr_ds, epochs=30,
                    #steps_per_epoch=int(np.ceil(tr_images.shape[0]/BATCH_SIZE)),
                    validation_data=val_ds,
                    #validation_steps=int(np.ceil(val_images.shape[0]/BATCH_SIZE)),
                    callbacks=[r_lr_cb, e_ly_cb]
                    )
```

```
WARNING:absl:lr is deprecated in Keras optimizer, please use learning_rate or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.A
Epoch 1/30
625/625 [=====] - 183s 270ms/step - loss: 1.8486 - accuracy: 0.3556 - val_loss: 1.7444 - val_accuracy: 0.3856 - lr:
Epoch 2/30
625/625 [=====] - 168s 269ms/step - loss: 1.4867 - accuracy: 0.4767 - val_loss: 1.4347 - val_accuracy: 0.4861 - lr:
Epoch 3/30
625/625 [=====] - 166s 266ms/step - loss: 1.3162 - accuracy: 0.5385 - val_loss: 2.1401 - val_accuracy: 0.3861 - lr:
Epoch 4/30
625/625 [=====] - 170s 272ms/step - loss: 1.1886 - accuracy: 0.5853 - val_loss: 1.2302 - val_accuracy: 0.5669 - lr:
Epoch 5/30
625/625 [=====] - 166s 265ms/step - loss: 1.0976 - accuracy: 0.6151 - val_loss: 1.4439 - val_accuracy: 0.5163 - lr:
Epoch 6/30
625/625 [=====] - 166s 265ms/step - loss: 1.0354 - accuracy: 0.6405 - val_loss: 1.0721 - val_accuracy: 0.6252 - lr:
Epoch 7/30
625/625 [=====] - 166s 265ms/step - loss: 0.9723 - accuracy: 0.6631 - val_loss: 1.6600 - val_accuracy: 0.4167 - lr:
Epoch 8/30
625/625 [=====] - 165s 264ms/step - loss: 0.9124 - accuracy: 0.6829 - val_loss: 1.1722 - val_accuracy: 0.5945 - lr:
Epoch 9/30
625/625 [=====] - 165s 264ms/step - loss: 0.8644 - accuracy: 0.6992 - val_loss: 1.0266 - val_accuracy: 0.6457 - lr:
Epoch 10/30
625/625 [=====] - 166s 265ms/step - loss: 0.8295 - accuracy: 0.7125 - val_loss: 1.5932 - val_accuracy: 0.4846 - lr:
Epoch 11/30
625/625 [=====] - 165s 264ms/step - loss: 0.7838 - accuracy: 0.7268 - val_loss: 1.9405 - val_accuracy: 0.4983 - lr:
Epoch 12/30
625/625 [=====] - 165s 265ms/step - loss: 0.7484 - accuracy: 0.7399 - val_loss: 1.2085 - val_accuracy: 0.5760 - lr:
Epoch 13/30
625/625 [=====] - 166s 265ms/step - loss: 0.7098 - accuracy: 0.7532 - val_loss: 1.1009 - val_accuracy: 0.6251 - lr:
Epoch 14/30
625/625 [=====] - ETA: 0s - loss: 0.6782 - accuracy: 0.7651
Epoch 14: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
625/625 [=====] - 166s 265ms/step - loss: 0.6782 - accuracy: 0.7651 - val_loss: 1.0624 - val_accuracy: 0.6346 - lr:
Epoch 15/30
625/625 [=====] - 166s 265ms/step - loss: 0.4031 - accuracy: 0.8587 - val_loss: 0.8420 - val_accuracy: 0.7259 - lr:
Epoch 16/30
625/625 [=====] - 166s 265ms/step - loss: 0.3171 - accuracy: 0.8888 - val_loss: 0.8779 - val_accuracy: 0.7212 - lr:
Epoch 17/30
625/625 [=====] - 164s 263ms/step - loss: 0.2703 - accuracy: 0.9054 - val_loss: 0.9732 - val_accuracy: 0.7121 - lr:
Epoch 18/30
625/625 [=====] - 164s 263ms/step - loss: 0.2401 - accuracy: 0.9157 - val_loss: 0.9357 - val_accuracy: 0.7170 - lr:
Epoch 19/30
625/625 [=====] - 165s 264ms/step - loss: 0.2108 - accuracy: 0.9277 - val_loss: 1.0204 - val_accuracy: 0.7137 - lr:
Epoch 20/30
625/625 [=====] - ETA: 0s - loss: 0.1893 - accuracy: 0.9337
Epoch 20: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.
625/625 [=====] - 165s 263ms/step - loss: 0.1893 - accuracy: 0.9337 - val_loss: 0.9874 - val_accuracy: 0.7211 - lr:
Epoch 21/30
625/625 [=====] - 164s 263ms/step - loss: 0.1440 - accuracy: 0.9513 - val_loss: 0.9368 - val_accuracy: 0.7377 - lr:
Epoch 22/30
625/625 [=====] - 165s 264ms/step - loss: 0.1283 - accuracy: 0.9576 - val_loss: 0.9412 - val_accuracy: 0.7340 - lr:
Epoch 23/30
625/625 [=====] - 165s 264ms/step - loss: 0.1242 - accuracy: 0.9578 - val_loss: 0.9698 - val_accuracy: 0.7351 - lr:
Epoch 24/30
625/625 [=====] - 165s 264ms/step - loss: 0.1170 - accuracy: 0.9610 - val_loss: 0.9873 - val_accuracy: 0.7381 - lr:
Epoch 25/30
625/625 [=====] - ETA: 0s - loss: 0.1144 - accuracy: 0.9619
Epoch 25: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.
```

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

